Conference -

Source URL: https://hc2024.hotchips.org/program/conference/

HC2024 Sponsors

Rhodium:

Platinum:

Gold:

Silver:

Conference

Conference Day 1: Monday, August 26, 2024

Time (PDT)

Title

Presenters

7:45AM-9:15AM

Breakfast/Registration

☐

9:15AM-9:30AM

Welcome

☐

☐

General Chair Welcome

Ron Diamant, General Chair

☐

Progam Co-Chairs Welcome

Rob Aitken & Larry Yang, PC Co-Chairs

9:30AM-11:00AM

High-Performance Processors Part 1

Chair: Ian Bratt

☐

☐

Snapdragon X Elite Qualcomm Oryon CPU: Design & Architecture Overview

Gerard Williams, Qualcomm

☐

Lunar Lake: Powering the Next Generation of AI PCs

Arik Gihon, Intel

☐

IBM Next Generation Processor and AI Accelerator

Chris Berry, IBM

11:00AM-11:30AM

Coffee Break

☐

11:30AM-1:00PM

Specialized Processors

Chair: Renu Raman

☐

☐

Blackhole and TT-Metalium - The Standalone AI Computer and its Programming Model

Jasmina Vasiljevic & Davor Capalija, Tenstorrent

☐

SK Hynix AI-Specific Computing Memory Solution: From AiM device to Heterogeneous AiMX-xPU System for Comprehensive LLM Inference

Guhyun Kim, SK Hynix

□

Built for the Edge: The next generation Intel□□ Xeon 6 SoC

Praveen Mosur, Intel

1:00PM-2:15PM

Lunch (1 hr 15 min)

□

2:15PM-3:15PM

Keynote #1

Chair: Ralph Wittig

□

□

Predictable Scaling and Infrastructure

Trevor Cai, OpenAI

3:15PM-4:15PM

AI Processors Part 1

Chair: Pradeep Dubey

□

□

NVIDIA Blackwell Platform: Advancing Generative AI and Accelerated Computing

Ajay Tirumala & Raymond Wong, NVIDIA

□

SambaNova SN40L RDU: Breaking the Barrier of Trillion+ Parameter Scale Gen AI Computing

Raghu Prabhakar, SambaNova

4:15PM-4:45PM

Coffee Break (1/2 hr)

□

4:45PM-6:45PM

AI Processors Part 2

Chair: David Weaver

□

□

Intel Gaudi 3 AI Accelerator: Architected for Gen AI Training and Inference

Roman Kaplan, Intel

□

AMD InstinctTM MI300X Generative AI Accelerator and Platform Architecture

Alan Smith & Vamsi Krishna Alla, AMD

□

An AI Compute ASIC with Optical Attach to Enable Next Generation Scale-up Architectures

Manish Mehta, Broadcom

□

FuriosaAI RNGD: A Tensor Contraction Processor for Sustainable AI Computing

June Paik, Furiosa

6:45PM-8:30PM

Reception

□

Conference Day 2: Tuesday, August 27, 2024

Time (PDT)

Title

Presenters

7:45AM-8:30AM

Breakfast/Registration

□

8:30AM-9:00AM

Poster Lightning Session

□

□

Poster Lightning Talks

□

9:00AM-10:30AM

AI Processors Part 3

Chair: Yasuo Ishii

□

□

AMD Versal□□¢ AI Edge Series Gen 2 for Vision and Automotive

Tomai Knopp & Jeffrey Chu, AMD

□

Onyx: A Programmable Accelerator for Sparse Tensor Algebra

Kalhan Koul, Stanford

□

Next Gen MTIA - Meta□□□s Recommendation Inference Accelerator

Mahesh Maddury & Pankaj Kansal, Meta

10:30AM-11:00AM

Coffee Break

□

11:00AM-12:30PM

Networking Processors

Chair: Jae W. Lee

□

□

DOJO: An Exa-Scale Lossy AI Network using the Tesla Transport Protocol over Ethernet (TTPoE)

Eric Quinnell, Tesla

□

ACF-S: An 8-Tbit/s SuperNIC for High-Performance Data Movement in AI & Accelerated Compute Networks

Shrijeet Mukherjee & Thomas Norrie, Enfabrica

□

4 Tbit/s Optical Compute Interconnect Chiplet for XPU-to-XPU Connectivity

Saeed Fathololoumi, Intel

12:30PM-1:45PM

Lunch (1 hr 15 min)

□

1:45PM-2:45PM

Keynote #2

Chair: Ian Bratt

□

□

The Journey to Life with AI Pervasiveness

Victor Peng, AMD

2:45PM-3:45PM

High-Performance Processors Part 2

Chair: Nhon Quach

□

□

Wafer-Scale AI: Enabling Unprecedented AI Compute Performance

Sean Lie, Cerebras

□

XiangShan: An Open-Source Project for High-Performance RISC-V Processors Meeting Industrial-Grade Standards

Kaifan Wang, Chinese Academy of Sciences

3:45PM-4:15PM

Coffee Break (1/2 hr)

□

4:15PM-6:15PM

High-Performance Processors Part 3

Chair: Lingjie Xu

□

□

AmpereOne: Sustainable Computing for AI & Cloud Native Workloads

Matthew Erler, Ampere Computing

□

Inside MAIA 100

Sherry Xu & Chandru Ramakrishnan, Microsoft

□

AMD Next Generation □□□Zen 5□□□ Core

Brad Cohen & Mahesh Subramony, AMD

□

MN-Core 2: Second-generation processor of MN-Core architecture for AI and general-purpose HPC applications

Jun Makino, Preferred Networks

6:15PM-6:30PM

IEEE TCMM Awards

□

□

IEEE TCMM Awards

Gabriel Southern, TCMM Chair

6:30PM-6:45PM

Closing Remarks

□

□

Closing Remarks

Jan-Willem van de Waerdt, Vice Chair